Pandoro

GPU Infrastructure for ML Research Teams

🌱 Powered by Clean Energy

pandoro.today

The Problem

Research teams conducting ML experiments face an impossible situation:

- **Experiments that should take hours take months** without GPU acceleration
- Navigating **financial uncertainty** with usage-based pricing
- **Institutional red tape** and IT approval delays
- Proprietary infrastructure that **undermines reproducibility**

```
What you need:
✓ GPUs
✓ High RAM
✓ Substantial storage

What's unclear:
✗ How to access it affordably
✗ How to avoid queue delays
✗ How to ensure reproducibility
```

The Four Barriers Researchers Face

💰 Financial Barriers

- Usage-based pricing creates unpredictable costs
- Hidden data transfer and storage fees
- Hardware purchases hard to justify early-stage

🏛 Institutional Barriers

- IT approval processes delay acquisition
- Multi-week queue delays for shared resources
- Mandated environments vs. research needs

🔬 Reproducibility Barriers

- Cloud providers hide hardware specs
- Results can’t be validated without environment info
- Inconsistent hardware across institutions

⚙️ Operational Barriers

- Vendor lock-in with proprietary configs
- Data transfer fees block migration
- No clear path to scale

Our Solution

Dedicated consumer GPU systems at fixed monthly pricing
with complete hardware transparency

🖥️ Dedicated Access

No shared queues, no IT approval, immediate access

📋 Full Transparency

Complete specs disclosed for reproducible research

💵 Fixed Pricing

Run unlimited experiments without tracking usage

Why This Approach Works

Hyperscalers sell infrastructure. We provide access to computers.

You don't need deployment pipelines, auto-scaling groups, or infrastructure orchestration. You need to run ML experiments on reliable hardware with known specifications.

Fixed pricing solves the prediction problem.

AWS, Google Cloud, and Azure offer reserved instances requiring accurate future capacity predictions. You cannot predict experiment duration when exploring new methodologies.

Research needs transparency, not abstraction.

Hyperscalers abstract hardware behind instance types and virtualization layers. We provide complete hardware specifications because scientific publication requires reproducible computational environments.

Clean Energy Computing

🌲 Pacific Northwest Renewable Energy

Our infrastructure runs on Washington state's electrical grid, one of the cleanest in the United States.

- **~90% hydroelectric and renewable**
- Minimize environmental impact of compute-intensive research
- Sustainability without premium pricing

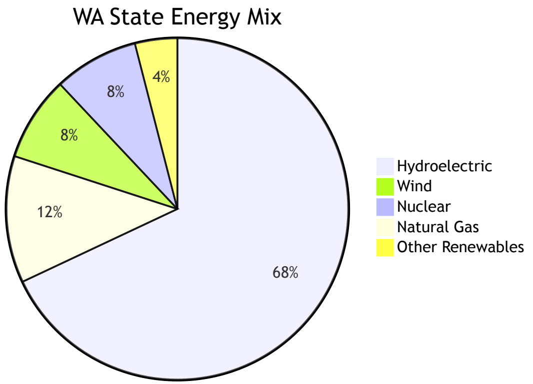

How Pandoro Compares

	Cloud Providers	University Computing	Hardware Purchase	**Pandoro**
Cost Model	Usage-based, unpredictable	Free but limited	High upfront	Fixed monthly
Access	Immediate	Weeks of queue delay	IT approval needed	Immediate
Hardware Specs	Hidden/abstracted	Often outdated	Known	Fully disclosed
Reproducibility	Difficult	Inconsistent	Good	Publication-ready
Migration Path	Vendor lock-in	N/A	N/A	Easy onsite transition

Who We Support

🔬 Domain Scientists

- Visual imaging ML projects
- Teams without ML engineering backgrounds
- Need hardware transparency for publishable results

🤖 Robotics Teams

- Computer vision experiments
- Sensor fusion and autonomous systems
- Need reliable compute without cloud unpredictability

Common Constraints

- Small teams without capital for hardware
- Institutional IT barriers
- Need for reproducible environments
- Must scale with project growth

Onsite Migration Path

Consumer-grade GPU systems enable easy transition to in-house infrastructure

No vendor lock-in. No workflow rewrites. No proprietary configurations.

About Bread Board Foundry

Pandoro is developed by **Bread Board Foundry** – we build specialized tools for teams working on meaningful, impactful projects.

Our Products

- **Pretzel** – Manufacturing optimization for hardware teams
- **Soufflé** – Goal-based planning for hardware leaders
- **Pandoro** – GPU compute for ML research

"Forging Software That Speaks Hardware"

breadboardfoundry.com

Let’s Connect

Schedule a Call

Discuss your research needs, timeline constraints, and budget considerations.

calendly.com/breadboardfoundry/pandoro

Learn More

Full hardware details, pricing structure, and technical specifications.

pandoro.today

Keenan Johnson • Bread Board Foundry • hello@breadboardfoundry.com

Questions?

🙋‍♀️ 🙋 🙋‍♂️

Thank you to Pascal Joly and the Sustainable AI community!